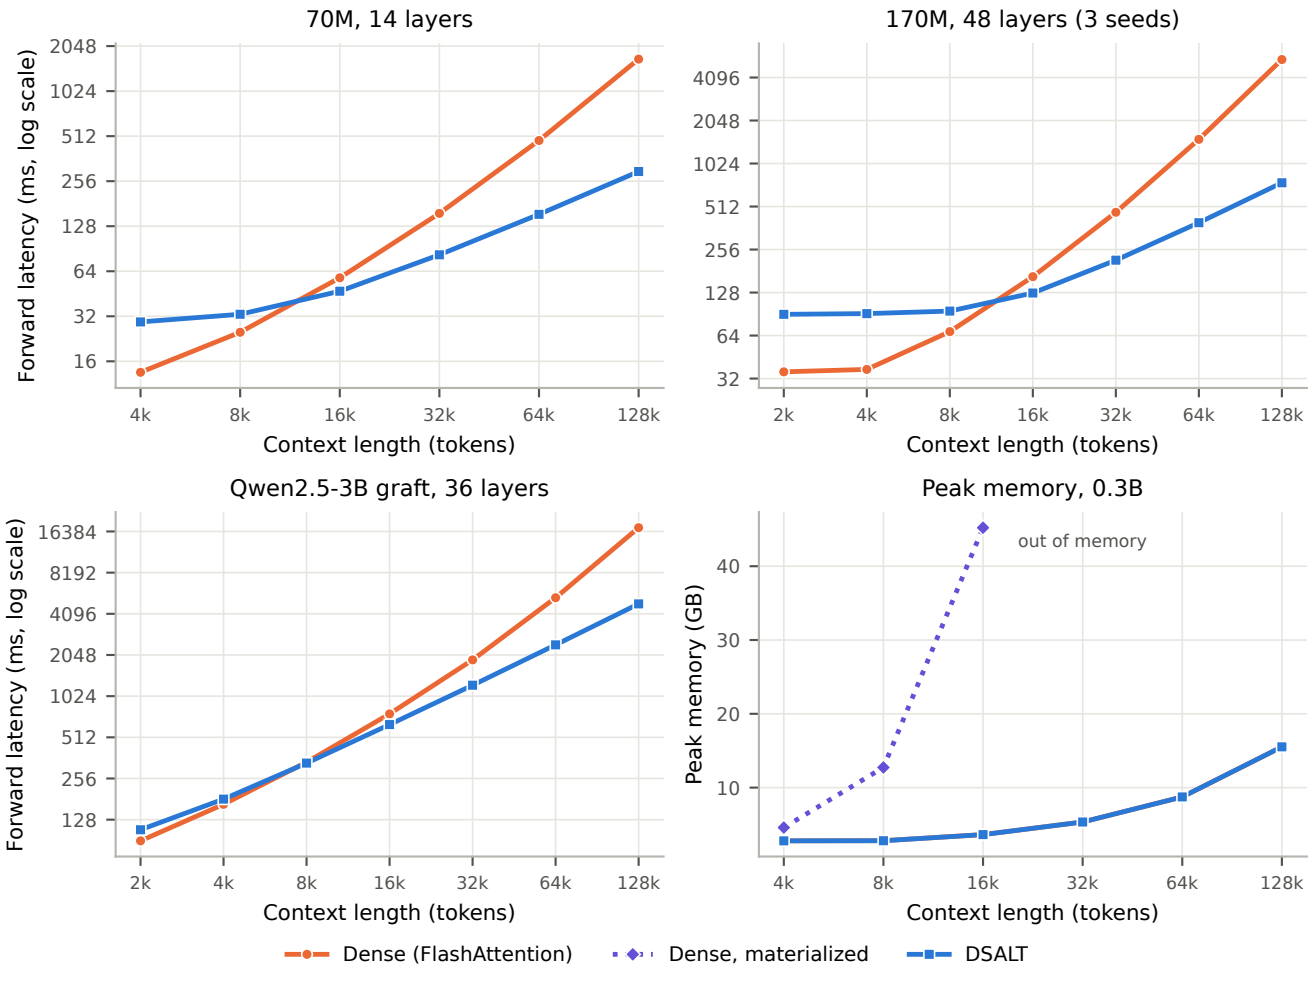
70M, 14 layers
Forward latency (ms, log scale)
2048
1024
512
256
128
64
32
16
4k
8k
16k
32k
64k
128k
Context length (tokens)
170M, 48 layers (3 seeds)
4096
2048
1024
512
256
128
64
32
2k
4k
8k
16k
32k
64k
128k
Context length (tokens)
Qwen2.5-3B graft, 36 layers
Forward latency (ms, log scale)
16384
8192
4096
2048
1024
512
256
128
2k
4k
8k
16k
32k
64k
128k
Context length (tokens)
Peak memory, 0.3B
Peak memory (GB)
out of memory
40
30
20
10
4k
8k
16k
32k
64k
128k
Context length (tokens)
Dense (FlashAttention)
Dense, materialized
DSALT